

# Statistical Language Modelling

## COM6513 Natural Language Processing

Nikos Aletras

Computer Science Department

Week 4  
Spring 2023



The  
University  
Of  
Sheffield.

In the previous lecture...

- Our first NLP problem: **Text classification**

## In the previous lecture...

- Our first NLP problem: **Text classification**
- But we ignored **word order** (apart from short sequences, e.g. n-grams)!

In this lecture...

## Our second NLP problem: Language Modelling

What is the probability of a given sequence of words in a particular language (e.g. English)?

In this lecture...

## Our second NLP problem: Language Modelling

What is the probability of a given sequence of words in a particular language (e.g. English)?

Odd problem. Applications?

# Applications of LMs

- Word likelihood for query completion in information retrieval  
(“Is Sheffield” → try it on your search engine)

# Applications of LMs

- Word likelihood for query completion in information retrieval (“Is Sheffield” → try it on your search engine)
- Language detection (“Ciao Sheffield” is it Italian or English?)

# Applications of LMs

- Word likelihood for query completion in information retrieval (“Is Sheffield” → try it on your search engine)
- Language detection (“Ciao Sheffield” is it Italian or English?)
- Grammatical error detection (“You’re welcome” or “Your welcome”?)



# Applications of LMs

- Word likelihood for query completion in information retrieval (“Is Sheffield” → try it on your search engine)
- Language detection (“Ciao Sheffield” is it Italian or English?)
- Grammatical error detection (“You’re welcome” or “Your welcome”?)
- Speech recognition (“I was tired too.” or “I was tired two.”?)

# Problem setup

Training data is a (often large) set of sentences  $\mathbf{x}^m$  with words  $x_n$ :

$$D_{train} = \{\mathbf{x}^1, \dots, \mathbf{x}^M\}$$
$$\mathbf{x} = [x_1, \dots, x_N]$$

for example:

$$\mathbf{x} = [\langle s \rangle, \text{The, water, is, clear, .., } \langle /s \rangle]$$

$\langle s \rangle$ : denotes start of the sentence

$\langle /s \rangle$ : denotes end of the sentence

# Calculate sentence probabilities

We want to learn a model that returns the **probability of an unseen sentence  $\mathbf{x}$** :

$$P(\mathbf{x}) = P(x_1, \dots, x_n), \text{ for } \forall \mathbf{x} \in V^{\max N}$$

$V$  is the vocabulary and  $V^{\max N}$  all possible sentences.

# Calculate sentence probabilities

We want to learn a model that returns the **probability of an unseen sentence  $\mathbf{x}$** :

$$P(\mathbf{x}) = P(x_1, \dots, x_n), \text{ for } \forall \mathbf{x} \in V^{maxN}$$

$V$  is the vocabulary and  $V^{maxN}$  all possible sentences.

How to compute probability?

# Unigram language model

Multiply the probability of each word appearing in the sentence  $\mathbf{x}$  computed over the entire corpus:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n) = \prod_{n=1}^N \frac{c(x_n)}{\sum_{x \in V} c(x)}$$

# Unigram language model

Multiply the probability of each word appearing in the sentence  $\mathbf{x}$  computed over the entire corpus:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n) = \prod_{n=1}^N \frac{c(x_n)}{\sum_{x \in V} c(x)}$$

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

# Unigram language model

Multiply the probability of each word appearing in the sentence  $\mathbf{x}$  computed over the entire corpus:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n) = \prod_{n=1}^N \frac{c(x_n)}{\sum_{x \in V} c(x)}$$

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

$$P(\text{i love}) = P(\text{i})P(\text{love}) = \frac{2}{20} \cdot \frac{1}{20} = 0.005$$

# What could go wrong?

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

- The most probable word is <s> ( $\frac{3}{20}$ )



# What could go wrong?

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

- The most probable word is <s> ( $\frac{3}{20}$ )
- The most probable single-word sentence is “<s>”

# What could go wrong?

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

- The most probable word is <s> ( $\frac{3}{20}$ )
- The most probable single-word sentence is "<s>"
- The most probable two-word sentence is "<s> <s>"

# What could go wrong?

<s> i love playing basketball </s>

<s> arctic monkeys are from sheffield </s>

<s> i study in sheffield uni </s>

- The most probable word is <s> ( $\frac{3}{20}$ )
- The most probable single-word sentence is "<s>"
- The most probable two-word sentence is "<s> <s>"
- The most probable  $N$ -word sentence is  $N \times$  "<s>"

# Maximum Likelihood Estimation

Instead of assuming **independence**:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n)$$

We assume that each word is **dependent** on all previous ones:

$$\begin{aligned} P(\mathbf{x}) &= P(x_1, \dots, x_N) \\ &= P(x_1)P(x_2 \dots x_N | x_1) \\ &= P(x_1)P(x_2 | x_1) \dots P(x_N | x_1, \dots, x_{N-1}) \\ &= \prod_{n=1}^N P(x_n | x_1, \dots, x_{n-1}) \quad (\text{chain rule}) \end{aligned}$$

# Maximum Likelihood Estimation

Instead of assuming **independence**:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n)$$

We assume that each word is **dependent** on all previous ones:

$$\begin{aligned} P(\mathbf{x}) &= P(x_1, \dots, x_N) \\ &= P(x_1)P(x_2 \dots x_N | x_1) \\ &= P(x_1)P(x_2 | x_1) \dots P(x_N | x_1, \dots, x_{N-1}) \\ &= \prod_{n=1}^N P(x_n | x_1, \dots, x_{n-1}) \quad (\text{chain rule}) \end{aligned}$$

What could go wrong?

# Problems with MLE

Let's analyse this:

$$P(\mathbf{x}) = P(x_1)P(x_2|x_1)P(x_3|x_2, x_1)\dots P(x_N|x_1, \dots, x_{N-1})$$

$$P(x_n|x_{n-1}\dots x_1) = \frac{c(x_1\dots x_{n-1}, x_n)}{c(x_1\dots x_{n-1})}$$

As we condition on more words, the counts become **sparser**

# Bigram Language Models

Assume that the choice of a word **depends only on the one before it**:

$$P(\mathbf{x}) = \prod_{n=1}^N P(x_n | x_{n-1}) = \prod_{n=1}^N \frac{c(x_{n-1}, x_n)}{c(x_{n-1})}$$

k-th order **Markov assumption**:

$$P(x_n | x_{n-1}, \dots, x_1) \approx P(x_n | x_{n-1}, \dots, x_{n-k})$$

with  $k=1$



# Bigram LM: From counts to probabilities

Unigram counts:

| <b>arctic</b> | <b>monkeys</b> | <b>are</b> | <b>my</b> | <b>favourite</b> | <b>band</b> |
|---------------|----------------|------------|-----------|------------------|-------------|
| 100           | 600            | 4000       | 3000      | 500              | 200         |

Bigram counts (rows:  $x_{i-1}$ , cols:  $x_i$ ):

|                  | <b>arctic</b> | <b>monkeys</b> | <b>are</b> | <b>my</b> | <b>favourite</b> | <b>band</b> |
|------------------|---------------|----------------|------------|-----------|------------------|-------------|
| <b>arctic</b>    | 0             | 10             | 2          | 0         | 0                | 0           |
| <b>monkeys</b>   | 0             | 0              | 250        | 1         | 5                | 0           |
| <b>are</b>       | 3             | 45             | 0          | 600       | 25               | 1           |
| <b>my</b>        | 0             | 2              | 0          | 1         | 300              | 5           |
| <b>favourite</b> | 0             | 1              | 0          | 0         | 0                | 50          |
| <b>band</b>      | 0             | 0              | 3          | 10        | 0                | 0           |



# Bigram LM: From counts to probabilities

From the bigram count matrix, compute probabilities by dividing each cell by the appropriate unigram count for its row.

Bigram probabilities (rows:  $x_{i-1}$ , cols:  $x_i$ ):

|                  | <b>arctic</b> | <b>monkeys</b> | <b>are</b> | <b>my</b> | <b>favourite</b> | <b>band</b> |
|------------------|---------------|----------------|------------|-----------|------------------|-------------|
| <b>arctic</b>    | 0             | 0.1            | 0.02       | 0         | 0                | 0           |
| <b>monkeys</b>   | 0             | 0              | 0.417      | 0.0017    | 0.008            | 0           |
| <b>are</b>       | 0.0008        | 0.0113         | 0          | 0.15      | 0.0063           | 0.00003     |
| <b>my</b>        | 0             | 0.0007         | 0          | 0.0003    | 0.1              | 0.0017      |
| <b>favourite</b> | 0             | 0.002          | 0          | 0         | 0                | 0.1         |
| <b>band</b>      | 0             | 0              | 0.015      | 0.05      | 0                | 0           |

## Example: Bigram language model

$\mathbf{x} = [\text{arctic}, \text{monkeys}, \text{are}, \text{my}, \text{favourite}, \text{band}]$

$$\begin{aligned} P(\mathbf{x}) &= P(\text{monkeys}|\text{arctic})P(\text{are}|\text{monkeys})P(\text{my}|\text{are}) \\ &\quad P(\text{favourite}|\text{my})P(\text{band}|\text{favourite}) \\ &= \frac{c(\text{arctic}, \text{monkeys})}{c(\text{arctic})} \cdots \frac{c(\text{favourite}, \text{band})}{c(\text{favourite})} \\ &= 0.1 \cdot 0.417 \cdot 0.15 \cdot 0.1 \cdot 0.1 \\ &= 0.00006255 \end{aligned}$$

# Longer contexts (N-gram LMs)

$$P(x|context) = \frac{P(context, x)}{P(context)} = \frac{c(context, x)}{c(context)}$$

- In bigram LM *context* is  $x_{n-1}$ , trigram  $x_{n-2}, x_{n-1}$ , etc.
- The longer the context:
  - the more likely to capture long-range dependencies:  
“I saw a tiger that was really very...” fierce or talkative?
  - the sparser the counts (zero probabilities)
- 5-grams and training sets with billions of tokens are common.

# Unknown Words

- If a word was never encountered in training, any sentence containing it will have probability 0
- It happens:
  - all corpora are finite
  - new words emerging

# Unknown Words

- If a word was never encountered in training, any sentence containing it will have probability 0
- It happens:
  - all corpora are finite
  - new words emerging
- Common solutions:
  - Generate unknown words in the training data by replacing low-frequency words with a special UNKNOWN token
  - Use classes of unknown words, e.g. names and numbers

# Implementation details

- Dealing with large datasets requires efficiency:
  - use log probabilities to avoid underflows (small numbers)
  - efficient data structures for sparse counts, e.g. lossy data structures Bloom filters)

# Implementation details

- Dealing with large datasets requires efficiency:
  - use log probabilities to avoid underflows (small numbers)
  - efficient data structures for sparse counts, e.g. lossy data structures Bloom filters)

How do we train and evaluate our language models?

- We need train/dev/test data
- Evaluation approaches

# Intrinsic Evaluation: Accuracy

- How well does our LM predict the next word?
- **I always order pizza with cheese and...**
  - mushrooms?
  - bread?
  - and?
- Accuracy: how often the LM predicts the correct word
- The higher the better



# Intrinsic Evaluation: Perplexity

- **Perplexity**: the inverse probability of the test set  $\mathbf{x} = [x_1, \dots, x_N]$ , normalised by the number of words  $N$ :

$$\begin{aligned} PP(\mathbf{x}) &= P(x_1, \dots, x_N)^{1/N} \\ &= \sqrt[N]{\frac{1}{P(x_1, \dots, x_N)}} \\ &= \sqrt[N]{\frac{1}{\prod_{i=1}^N P(x_i | x_1, \dots, x_{i-1})}} \end{aligned}$$

# Intrinsic Evaluation: Perplexity

- **Perplexity:** the inverse probability of the test set  $\mathbf{x} = [x_1, \dots, x_N]$ , normalised by the number of words  $N$ :

$$\begin{aligned} PP(\mathbf{x}) &= P(x_1, \dots, x_N)^{1/N} \\ &= \sqrt[N]{\frac{1}{P(x_1, \dots, x_N)}} \\ &= \sqrt[N]{\frac{1}{\prod_{i=1}^N P(x_i | x_1, \dots, x_{i-1})}} \end{aligned}$$

- Measures how well a probability distribution predicts a sample.
- The lower the better.

Why is a bigram language model likely to have lower perplexity than a unigram one?

# Intrinsic Evaluation: Perplexity

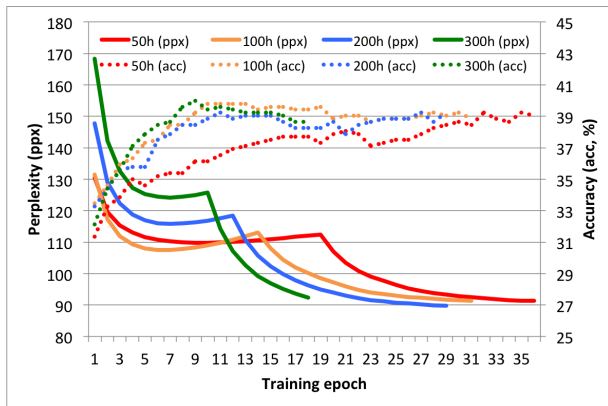
- **Perplexity:** the inverse probability of the test set  $\mathbf{x} = [x_1, \dots, x_N]$ , normalised by the number of words  $N$ :

$$\begin{aligned} PP(\mathbf{x}) &= P(x_1, \dots, x_N)^{1/N} \\ &= \sqrt[N]{\frac{1}{P(x_1, \dots, x_N)}} \\ &= \sqrt[N]{\frac{1}{\prod_{i=1}^N P(x_i | x_1, \dots, x_{i-1})}} \end{aligned}$$

- Measures how well a probability distribution predicts a sample.
- The lower the better.

Why is a bigram language model likely to have lower perplexity than a unigram one? **There is more context to predict the next word!**

# The problem with perplexity



- Doesn't always correlate with application performance
- Can't evaluate non probabilistic LMs

# Extrinsic Evaluation

- Sentence completion
- Grammatical error correction: detecting “odd” sentences and propose alternatives
- Natural language generation: prefer more “natural” sentences
- Speech recognition
- Machine translation

# Smoothing

- What happens when words that are in our vocabulary appear with an unseen context in test data?

# Smoothing

- What happens when words that are in our vocabulary appear with an unseen context in test data?
- They will be assigned with zero probability

# Smoothing

- What happens when words that are in our vocabulary appear with an unseen context in test data?
- They will be assigned with zero probability




**Smoothing (or discounting)** to the rescue: Steal from the rich and give to the poor!



# Smoothing intuition


- We often want to make estimates from sparse statistics:

$P(w \mid \text{denied the})$   
3 allegations  
2 reports  
1 claims  
1 request  
7 total



- Smoothing flattens spiky distributions so they generalize better

$P(w \mid \text{denied the})$   
2.5 allegations  
1.5 reports  
0.5 claims  
0.5 request  
2 other  
7 total



Taking from the frequent and giving to the rare (discounting)

# Add-1 Smoothing

Add-1 (or Laplace) smoothing adds one to all bigram counts:

$$P_{add-1}(x_n|x_{n-1}) = \frac{c(x_{n-1}, x_n) + 1}{c(x_{n-1}) + |V|}$$

Pretend we have seen all bigrams at least once!

# Add-k Smoothing

Add-1 puts too much probability mass to unseen bigrams, better to add- $k$ ,  $k < 1$ :

$$P_{add-k}(x_n|x_{n-1}) = \frac{counts(x_{n-1}, x_n) + k}{counts(x_{n-1}) + k|V|}$$

# Add-k Smoothing

Add-1 puts too much probability mass to unseen bigrams, better to add- $k$ ,  $k < 1$ :

$$P_{add-k}(x_n|x_{n-1}) = \frac{counts(x_{n-1}, x_n) + k}{counts(x_{n-1}) + k|V|}$$

$k$  is a hyperparameter: choose optimal value on the dev set!

# Interpolation

Longer contexts are more informative:

dog bites ... better than bites ...

but only if they are frequent enough:

canid bites ...better than bites ...?

# Interpolation

Longer contexts are more informative:

dog bites ... better than bites ...

but only if they are frequent enough:

canid bites ...better than bites ...?

Can we combine evidence from unigram, bigram and trigram probabilities?

# Simple Linear Interpolation

For a trigram LM:

$$\begin{aligned} P_{SLI}(x_n|x_{n-1}, x_{n-2}) = & \lambda_3 P(x_n|x_{n-1}, x_{n-2}) \\ & + \lambda_2 P(x_n|x_{n-1}) \\ & + \lambda_1 P(x_n) \quad \lambda_i > 0, \sum \lambda_i = 1 \end{aligned}$$

# Simple Linear Interpolation

For a trigram LM:

$$\begin{aligned} P_{SLI}(x_n|x_{n-1}, x_{n-2}) = & \lambda_3 P(x_n|x_{n-1}, x_{n-2}) \\ & + \lambda_2 P(x_n|x_{n-1}) \\ & + \lambda_1 P(x_n) \quad \lambda_i > 0, \sum \lambda_i = 1 \end{aligned}$$

- Weighted average of unigram, bigram and trigram probabilities



# Simple Linear Interpolation

For a trigram LM:

$$\begin{aligned} P_{SLI}(x_n|x_{n-1}, x_{n-2}) = & \lambda_3 P(x_n|x_{n-1}, x_{n-2}) \\ & + \lambda_2 P(x_n|x_{n-1}) \\ & + \lambda_1 P(x_n) \quad \lambda_i > 0, \sum \lambda_i = 1 \end{aligned}$$

- Weighted average of unigram, bigram and trigram probabilities
- How we choose the value of  $\lambda$ s?

# Simple Linear Interpolation

For a trigram LM:

$$\begin{aligned} P_{SLI}(x_n|x_{n-1}, x_{n-2}) = & \lambda_3 P(x_n|x_{n-1}, x_{n-2}) \\ & + \lambda_2 P(x_n|x_{n-1}) \\ & + \lambda_1 P(x_n) \quad \lambda_i > 0, \sum \lambda_i = 1 \end{aligned}$$

- Weighted average of unigram, bigram and trigram probabilities
- How we choose the value of  $\lambda$ s? Parameter tuning on the dev set!

# Backoff

Start with n-gram order of  $k$  but if the counts are 0 use  $k - 1$ :

$$BO(x_n|x_{n-1} \dots x_{n-k}) = \begin{cases} P(x_n|x_{n-1} \dots x_{n-k}), & \text{if } c(x_n \dots x_{n-k}) > 0 \\ BO(x_n|x_{n-1} \dots x_{n-k+1}), & \text{otherwise} \end{cases}$$

# Backoff

Start with n-gram order of  $k$  but if the counts are 0 use  $k - 1$ :

$$BO(x_n | x_{n-1} \dots x_{n-k}) = \begin{cases} P(x_n | x_{n-1} \dots x_{n-k}), & \text{if } c(x_n \dots x_{n-k}) > 0 \\ BO(x_n | x_{n-1} \dots x_{n-k+1}), & \text{otherwise} \end{cases}$$

Is this a probability distribution?

**NO!** Must discount probabilities for contexts with counts  $P^*$  and distribute the mass to the shorter context ones:

$$P_{BO}(x_n | x_{n-1} \dots x_{n-k}) = \begin{cases} P^*(x_n | x_{n-1} \dots x_{n-k}), & \text{if } c(x_n \dots x_{n-k}) > 0 \\ \alpha^{x_{n-1} \dots x_{n-k}} P_{BO}(x_n | x_{n-1} \dots x_{n-k+1}), & \text{otherwise} \end{cases}$$

$$\alpha^{x_{n-1} \dots x_{n-k}} = \frac{\beta^{x_{n-1} \dots x_{n-k}}}{\sum P_{BO}(x_n | x_{n-1} \dots x_{n-k+1})}$$

$\beta$ , is the left-over probability mass for the (n-k)-gram

# Absolute Discounting

Using 22M words for train and held-out

| Bigram count in training | Bigram count in heldout set |
|--------------------------|-----------------------------|
| 0                        | .0000270                    |
| 1                        | 0.448                       |
| 2                        | 1.25                        |
| 3                        | 2.24                        |
| 4                        | 3.23                        |
| 5                        | 4.21                        |
| 6                        | 5.23                        |
| 7                        | 6.21                        |
| 8                        | 7.21                        |
| 9                        | 8.26                        |

Can you predict the heldout (test) set average count given the training?

# Absolute Discounting

Using 22M words for train and held-out

| Bigram count in training | Bigram count in heldout set |
|--------------------------|-----------------------------|
| 0                        | .0000270                    |
| 1                        | 0.448                       |
| 2                        | 1.25                        |
| 3                        | 2.24                        |
| 4                        | 3.23                        |
| 5                        | 4.21                        |
| 6                        | 5.23                        |
| 7                        | 6.21                        |
| 8                        | 7.21                        |
| 9                        | 8.26                        |

Can you predict the heldout (test) set average count given the training?

Testing counts = training counts - 0.75 (absolute discount)

# Absolute discounting

$$P_{AbsDiscount}(x_n|x_{n-1}) = \frac{c(x_n, x_{n-1}) - d}{c(x_{n-1})} + \lambda_{x_{n-1}} P(x_n)$$

- $d = 0.75$ ,  $\lambda$ s tuned to ensure we have a valid probability distribution.
- Component of the **Kneser-Ney** discounting:
  - Intuition: a word can be very frequent, but if only follows very few contexts,  
e.g. **Francisco** is frequent but almost always follows **San**
  - The unigram probability in the context of the bigram should capture how likely  $x_n$  is to be a novel continuation.



# Stupid Backoff

- Do we really need probabilities? Estimating the additional parameters takes time for large corpora.

# Stupid Backoff

- Do we really need probabilities? Estimating the additional parameters takes time for large corpora.
- If scoring is enough, **stupid backoff** works adequately:

$$SBO(x_n | x_{n-1} \dots x_{n-k}) = \begin{cases} P(x_n | x_{n-1} \dots x_{n-k}), & \text{if } c(x_n \dots x_{n-k}) > 0 \\ \lambda SBO(x_n | x_{n-1} \dots x_{n-k+1}), & \text{otherwise} \end{cases}$$

- Empirically found that  $\lambda = 0.4$  works well

# Stupid Backoff

- Do we really need probabilities? Estimating the additional parameters takes time for large corpora.
- If scoring is enough, **stupid backoff** works adequately:

$$SBO(x_n | x_{n-1} \dots x_{n-k}) = \begin{cases} P(x_n | x_{n-1} \dots x_{n-k}), & \text{if } c(x_n \dots x_{n-k}) > 0 \\ \lambda SBO(x_n | x_{n-1} \dots x_{n-k+1}), & \text{otherwise} \end{cases}$$

- Empirically found that  $\lambda = 0.4$  works well
- They called it stupid because they didn't expect it to work well!

# Last words: More data defeats smarter models!




Figure 5: BLEU scores for varying amounts of data using Kneser-Ney (KN) and Stupid Backoff (SB).

From [Large Language Models in Machine Translation](#)

# Bibliography

- Chapter 3 from Jurafsky & Martin
- Chapter 6 from Eisentein
- Michael Collins' [notes](#) on LMs

# Coming up next...

- We have learned how to model word sequences using Markov models
- In the following lecture we will look at how to perform part-of-speech tagging using:
  - the **Hidden** Markov Model (HMM)
  - the **Conditional Random Fields** (CRFs), an extension of logistic regression for sequence modelling